Student Personal AI Assistant: A Strategic Build-and-Scale Guide (India-First, 2026)

TL;DR

- **The problem is real but the wedge is wrong.** Academic information overload and missed deadlines are genuine pains for Indian students, but the founder's instinct to build "55 features" across Gmail + Classroom + Calendar + burnout-scoring is a trap. The single sharpest, cheapest, most-defensible first feature is a **"What's due this week?" deadline radar built on Google Classroom + Calendar data — deliberately avoiding Gmail at first** — because Classroom/Calendar are merely "sensitive" OAuth scopes while Gmail is "restricted" and triggers an expensive annual CASA security assessment.
- **Pathway is not needed.** At pilot and early scale, event-driven webhooks (Gmail push via Pub/Sub) + scheduled polling (Classroom/Calendar) + Postgres/pgvector on Supabase is sufficient, far cheaper, and easier to operate. LLM cost per active student is near-zero (~₹2–5/month) using Llama 3.1-8B on Groq for text and Gemini Flash-Lite for multimodal — so the binding constraints are OAuth verification, distribution, and retention, not compute.
- **"Untouchable for 5 years" is not realistic; a 2–3 year India-campus head-start is.** No solo founder can out-engineer Google. The only durable moat is **deep, per-college integration + proprietary deadline-extraction feedback data + campus community/ambassador lock-in**, college by college. Validate ruthlessly before scaling: if you cannot get 40%+ weekly retention in a single-campus pilot, do not proceed.

Key Findings

1. **Demand is real and painful, but "deadline-missing" needs first-party validation.** India recorded 14,488 student suicides in 2024 (8.5% of the national total, up 4.3%

from 2023) with academic pressure as a leading driver; broad academic-overload and procrastination statistics exist. But almost none of this directly measures "missed a deadline because info was buried in email/portals." The founder must generate that evidence at their own campus.

2. **The competitive field is crowded at the edges but empty in the middle.** Reclaim, Motion, Shortwave, Superhuman, Notion AI, and Google's own Gemini all solve adjacent problems for professionals — none is built for the *Indian college student's academic logistics* (Classroom assignments, college circulars, vernacular notices, mess menus). College ERPs (Samarth, 200+ HEIs) and edtech (PW, Unacademy) are content/admin plays, not personal-assistant plays. There is a genuine unoccupied niche.
3. **Google API compliance is the #1 real-world blocker — and there is a clean path around it.** Gmail's `gmail.readonly` is a *restricted* scope that forces an annual third-party CASA security assessment (~$540/yr minimum via TAC Security) if data touches your servers. Classroom and Calendar scopes are merely *sensitive* — verification only, no CASA. A <100-user pilot needs no verification at all.
4. **Compute is cheap; the model stack is a solved problem.** Llama 3.1-8B on Groq costs $0.05/$0.08 per million input/output tokens; Gemini 2.0/2.5 Flash-Lite costs $0.075–0.10/$0.30–0.40. Batching and prompt-caching cut this ~75%.
5. **Willingness to pay is very low.** Indian students anchor to Spotify Premium Student at ₹69/month (promo as low as ₹59). A productivity app's realistic ceiling is ₹49–99/month, with ~1–3% free-to-paid conversion. B2B2C (college/department licensing) is the more credible revenue path.

Details

1. Problem Validation & Demand

The macro evidence is strong on stress, weaker on the specific "logistics" pain. India recorded 14,488 student suicides in 2024 (7,669 male, 6,819 female; 8.5% of the national total of 1,70,746), a record high per NCRB's "Accidental Deaths and Suicides in India 2024" report, as reported by The Wire (11 May 2026). Academic pressure is repeatedly cited as a leading driver; one widely circulated figure is that 61% of suicide attempts at IITs were linked to academic stress. Nearly half of college students report academic overload and procrastination in US studies (~44.5%), and email-overload research finds professionals lose ~10.8 hours/week to non-critical email and that 68% say email overload contributes to burnout. Around 65% of students report skipping an assignment at least once due to overwhelm.

The honest gap: none of this proves students miss deadlines *because* information is fragmented across Gmail/Classroom/WhatsApp/college portals. That is the founder's core hypothesis and it is *unvalidated*. The stress data tells you the emotional stakes are high (good for messaging); it does not tell you students will adopt a deadline tool.

Cheap validation playbook (4–6 weeks, ~₹0–5,000):

- **15–20 problem interviews** at IIT Dharwad using "The Mom Test" method — ask about past behavior ("walk me through the last deadline you missed"), never pitch. Decision gate: do ⩾60% describe a specific, recent, painful instance unprompted?
- **A "fake-door" landing page** (Carrd/Vercel) describing the deadline radar + a waitlist email capture. Push via campus WhatsApp/Instagram groups. Gate: ⩾100 signups and ⩾25% click on "how it works" from one campus.
- **A manual "concierge MVP":** for 10 volunteers, manually read their Classroom + circulars for two weeks and WhatsApp them a "what's due" digest each morning. This validates value with zero code. Gate: do ⩾5 of 10 say they'd be "very disappointed" without it (Sean Ellis PMF test)?
- **Validation metrics that matter:** % who return unprompted (D7/D30 retention), digest open rate, "very disappointed" score ⩾40%, and unsolicited referrals. Vanity metrics (signups, pageviews) do *not* count as validation.

2. Competitive Landscape

- **General AI/email/calendar assistants:** Reclaim.ai (free tier; paid from ~$8–10/mo; acquired by Dropbox, announced 20 August 2024 — at the time Reclaim served 320,000+ users across 43,000+ companies and had raised $9.5M; no native mobile app, PWA only), Motion (~$29/mo individual; no free plan), Shortwave (free → ~$7–30/mo), Superhuman (~$40/mo; acquired by Grammarly in 2025 — Grammarly's parent rebranded to "Superhuman" on 29 Oct 2025 and the app is now "Superhuman Mail"), Notion AI. All target *working professionals*, price in dollars, and ignore Classroom/college-circular data. Google's own Gemini in Workspace can summarize email and is the existential threat (see Moat).
- **Student productivity/study apps:** My Study Life, Todoist, Notion for students, academic planners — all *manual entry*. None auto-ingest the student's actual academic data sources. That manual-entry friction is precisely the wedge.
- **Indian edtech & ERP:** Physics Wallah (profitable, IPO'd, tier-2/3 vernacular strength), Unacademy (total revenue down 16% YoY to ₹826.3 Cr in FY25 from ₹988 Cr, net loss cut 31% to ₹436 Cr per a document reviewed by Entrackr; valuation fell from ~$3.5B in 2021 to under $500M, with upGrad signing a term sheet to acquire it in an all-stock deal in March 2026), BYJU'S (collapsed). These are *content* businesses. College ERPs like Samarth eGov (200+ HEIs, 40+ central/state universities, 100+ colleges) are *administrative* systems students are forced to use — clunky, not assistant-like.
- **The gap:** No one owns "academic logistics intelligence" — turning the messy stream of assignments, circulars, deadlines, and notices into a proactive personal assistant for the Indian student. **Why no one dominates it:** it's unglamorous, the data is fragmented and college-specific (hard to generalize), OAuth/compliance is a moat-in-reverse that scares hobbyists, and the willing-to-pay economics are poor — so big players ignore it and hobbyists can't get past verification.

3. Moat & Defensibility — A Candid Assessment

Blunt truth: "the #1 player, untouchable for 5 years" is not achievable for a solo student

founder against Google. Google can replicate deadline-extraction from Gmail/Classroom as a Gemini feature whenever it chooses. Anyone telling the founder otherwise is selling hype. What *is* achievable is a **2–3 year head start in a niche too small and unglamorous for Google to prioritize**, defended by moats Google structurally won't build:

- **Deep college-specific integration (strongest moat):** Google will never scrape *your specific college's* notice board, mess menu, exam circular PDFs, or department WhatsApp broadcasts. Per-college connectors and parsers are tedious, local, and exactly what a motivated student founder can do that a global platform won't.
- **Proprietary feedback loop:** every "is this actually a deadline? did the date parse correctly?" correction trains a private dataset of how Indian academic communications encode deadlines (Hinglish, "submit by EOD tomorrow," circular formats). This compounds and is not portable to a competitor.
- **Community/switching costs:** campus ambassador networks, club partnerships, and a per-college "my whole batch uses it" effect create local network density. Switching means re-onboarding your social graph.
- **Institutional partnerships:** a signed pilot with a department/college (even unpaid) is a relationship Google's self-serve model won't pursue and a competitor can't easily dislodge.

The strategy that survives the Google threat: stay narrow and hyper-local, win Indian campuses one at a time, accumulate per-college integration depth and correction data, and convert that into institutional relationships. Do **not** try to be a horizontal "AI assistant" — that's where Google wins. Be the thing that knows *IIT Dharwad's* academic rhythm better than any global tool ever will, then repeat.

4. Tech Architecture for Low Cost & Scale (2026)

(a) Do you need Pathway? No. Pathway is a real-time incremental-indexing streaming engine — genuinely useful only when a knowledge base changes constantly and freshness in seconds matters. A student deadline tool's data changes a few times a day. Use:

- **Gmail (later phase):** push notifications via `users.watch` → Google Cloud Pub/Sub → webhook, with `historyId` cursors for incremental sync. (Note: practitioners report Pub/Sub `watch` must be renewed at least every 7 days — recommend daily — and occasionally needs a polling fallback when notifications silently stop; budget for that reliability work.) Medium
- **Classroom/Calendar:** scheduled polling (cron, every few hours) + incremental sync. Simple and robust.
- **Storage/index:** Supabase (Postgres + pgvector, free tier → $25/mo Pro). Hybrid keyword + vector retrieval in one database. This replaces Pathway, Pinecone, and a separate vector DB. Metacto
- **Orchestration:** plain Python + a thin LangChain/LangGraph layer only where it earns its keep. Avoid heavy frameworks early.

(b) LLM strategy — cheapest viable in 2026:

- **Text (classification, extraction, Q&A):** Llama 3.1-8B Instant on Groq — $0.05/M input, $0.08/M output, ~560–840 tokens/sec. Use JSON mode for structured deadline extraction. Eesel AI
- **Multimodal (rubric/circular/mess-menu PDFs & images):** Gemini 2.5 Flash-Lite or 2.0 Flash-Lite — $0.075–0.10/M input, $0.30–0.40/M output, with a usable free tier. Price Per Token Price Per Token
- **Cost levers:** Groq Batch API (−50%) for nightly classification; prompt caching (−50%) for the repeated system prompt; together ~25% of on-demand. Provider risk note: ~90% of Groq's engineering staff (and founder Jonathan Ross) moved to NVIDIA in a Dec 2025 deal, though GroqCloud still operates independently — keep a provider-abstraction layer (the same Llama model is available via Together, Cerebras, DeepInfra at comparable or lower prices; DeepInfra lists 8B as low as ~$0.02–0.03/M). CloudZero

(c) Vector DB: pgvector inside Supabase (free, included) is sufficient to ~100K vectors. Qdrant/Chroma self-hosted or Pinecone free tier are alternatives, but a second datastore is unnecessary complexity early.

(d) Hosting on a low budget: Vercel (frontend) + Supabase (DB/auth/edge functions) + a small always-on worker on Fly.io/Railway/Render for the Pub/Sub webhook and cron. Cloudflare Workers + R2 for edge/static. Apply for credits: Microsoft for Startups Founders Hub (~$1–5K, no VC needed), Cloudflare for Startups ($5K+), Google for Startups Cloud Program (up to ~$200K+ if accepted), Supabase ($1–10K). Oracle Cloud Always Free as a fallback always-on VM. Caution: free tiers can have "cliffs" (e.g., MongoDB Atlas M0 connection limits, AWS 12-month expiry) — prefer tiers that grow with you (Cloudflare Workers/R2, Supabase, Neon). Startupbricks Blog + 3

(e) Mobile vs web: Indian students are mobile-first, but native apps are costly to build/maintain solo and the OAuth flow is identical on web. **Start with a mobile-optimized PWA** (installable, push notifications via web push) — Reclaim itself ships only a PWA. Add a thin native shell (Expo/React Native) only after retention is proven. Delivery of the daily digest via **WhatsApp** is likely more important than any app UI for India. Max Productive AI

(f) Per-user cost: A typical student receiving ~10 academic items/day generates ~300 classification calls/month (~600 tokens each) plus ~30 Q&A queries — roughly **$0.02–0.05 (₹2–4) per active user per month** in LLM cost even before batching/caching. **What drives cost as you scale:** (1) free-tier abuse / inactive accounts still being polled — gate polling on active users; (2) multimodal PDF volume — cache parsed results; (3) CASA + infra fixed costs once you cross 100 users and add Gmail; (4) WhatsApp Business API messaging fees if you use the official API at scale.

5. Google API / Compliance Reality (Critical)

This is the make-or-break operational constraint, and the research yields a clear, exploitable asymmetry confirmed against Google's official documentation:

- **Scope classification (decisive):** `gmail.readonly` and `gmail.metadata` are **RESTRICTED** scopes — explicitly on Google's restricted-scopes list (support.google.com/cloud/answer/13464325). Classroom scopes (`classroom.courses.readonly`, `classroom.coursework.me.readonly`, `classroom.announcements.readonly`) and `calendar.events` are merely **SENSITIVE** — they do *not* appear on the restricted list and do *not* trigger CASA. Google's own sensitive-scope documentation uses "reading events stored in Google Calendar" as its canonical example of a *sensitive* (not restricted) scope.
- **What "restricted" costs you:** if your app stores or transmits restricted-scope (Gmail) data on/through your servers, you must pass an **annual third-party CASA security assessment.** Tier 2 via TAC Security (Google's preferred, India-based lab) is **$540/app/year** under the basic plan (a discounted rate negotiated by Google), ~$720 for a "Premium" plan with unlimited rescans, and ~$1,800 for an "Enterprise Tier 2" plan; Tier 3 is ~$4,500. This is on top of OAuth verification (consent screen, privacy policy on a first-party domain, demo video, brand verification), which can take several weeks. Note Google's CASA team ceased the old PwC self-scan portal and now routes developers to paid assessors. Google Switch Labs
- **The 100-user cap:** an unverified app can have at most **100 users who grant the sensitive/restricted scopes, cumulative over the project's entire lifetime, non-resettable** (per support.google.com/cloud/answer/15549945 and the verification FAQ). Crucially, you must publish "In Production" (unverified) — "Testing" mode expires user authorizations/refresh tokens every 7 days and is unusable for a real pilot.
- **The CASA "no servers" exception is real but unsafe to rely on:** policy text ties the assessment to apps that "store or transmit restricted scope data on servers" / have "the capability to access Google user data from or through a server." A pure client-side/on-device architecture *could* therefore avoid CASA, but Google does not publish an explicit on-device exemption, the "capability to access via a server" language is a reviewer loophole, and there is a 2025 Google Cloud Community report of Google

denying the local-only carve-out for a Gmail iOS app. Verification is still required regardless. Do not bet the company on it. Google

The practical path (this is the strategic unlock):

1. **Pilot (<100 users): no verification, no CASA.** Run the home-campus pilot "In Production," unverified, under 100 lifetime users, indefinitely.
2. **Scale on SENSITIVE scopes only.** Build the core product on **Classroom + Calendar** — sensitive scopes that require OAuth verification but **no CASA security assessment**. This lets you scale to thousands of users for the cost of a verification review, not a $540+/yr audit.
3. **Add Gmail last, deliberately.** Only when Gmail-derived deadlines prove necessary, accept the restricted-scope CASA cost (~$540/yr) — by then you should have revenue or funding. Consider minimizing scope (metadata-only, or processing client-side) to reduce the assessment burden.

6. Monetization & Unit Economics

Willingness to pay is the hard ceiling. Indian students anchor to Spotify Premium Student at **₹69/month** per Spotify's official India student page (promotional rate as low as ₹59 after a free trial — roughly 50% off standard). A new productivity app cannot exceed that; realistic premium pricing is **₹49–99/month** or **₹399–699/year** (annual heavily preferred — students pay once per semester, not monthly). Expect **1–3% free-to-paid conversion**.

Recommended freemium split:

- **Free forever (core):** deadline radar for Classroom + Calendar, daily "what's due" digest (in-app + WhatsApp), natural-language Q&A, basic reminders. This must be genuinely useful — it's your growth engine.
- **Premium (₹49–99/mo or ₹499/yr):** Gmail integration (auto-extract deadlines from professor emails), multimodal PDF/image parsing (rubrics, circulars), the "academic

risk score"/overload predictor, smart adaptive reminders, multi-source merge (GitHub, Telegram, college portal), and priority/early features.

Unit economics (per active user/month): LLM/API cost ~₹2–5 (free tier, text-only) rising to ~₹10–20 for premium (multimodal + Gmail). Even at ₹49/mo, gross margin per paying user is ~90%+. **The problem is not margin — it's conversion volume.** At 1.5% conversion and ₹49/mo, 100,000 users → 1,500 payers → ~₹73,500/mo (~$880) gross revenue. That is a lifestyle-business trajectory on consumer pricing alone.

The better revenue path — B2B2C: license to **colleges/departments** (₹50,000-₹3,00,000/college/year for an institution-wide "student success / deadline-compliance" deployment), to **placement cells** (deadline-compliance + an internship/career feature), or **sponsorships** (ed-brands reaching engaged students). One signed college can exceed thousands of individual subscriptions and aligns with the institutional-funding path the founder already wants. This is also the most defensible revenue (relationship + switching cost). The edtech market lesson is unambiguous: Physics Wallah survived by building unit economics first; BYJU'S/Unacademy nearly died chasing growth — so prioritize sustainable economics over vanity scale.

7. Go-to-Market & Growth

- **0 → 100 (home campus, IIT Dharwad):** manual concierge + PWA. Recruit through your own hostel/department, clubs, and class WhatsApp groups. Hand-onboard every user. Goal: prove retention, not scale. (Tinder's playbook: it hit ~99.5% daily retention with a few hundred users *before* going to campuses — get retention right first.)
- **100 → 1,000 (saturate home campus + 1–2 nearby):** campus ambassador program (proven playbook: Tinder, Bumble, Red Bull globally, and in India OnePlus/Unstop/Google DSC/Gemini Student Ambassadors). Pay ambassadors in incentives/certificates (₹5,000-style response targets are standard) + a leaderboard. Partner with official student clubs and the placement cell. Seed Instagram/WhatsApp study communities.

- **1,000 → 100,000 (college-by-college expansion):** template the per-college integration (each new college = a connector config + a few circular-format parsers). Launch each campus via a local ambassador + a "your whole batch is on it" referral mechanic. Prioritize IITs/NITs first (homogeneous Google Workspace + Classroom usage, high digital literacy, tight networks), then private universities, then state universities (where vernacular/Hinglish support matters most).
- **Virality mechanics:** shareable "what's due this week" cards, "invite your batchmates to sync the same class" group value, and exam-season spikes. Distribution via WhatsApp is the single most important Indian growth channel.

8. Phased Roadmap

- **Phase 0 — Validation (Weeks 1–6; cost ~₹0–5K).** 15–20 Mom-Test interviews, fake-door landing page, 10-person concierge digest. **Gate to proceed:** ⩾40% "very disappointed" PMF score + ⩾5/10 concierge users retaining weekly. *If you fail this gate, stop or pivot — do not build.*
- **Phase 1 — MVP (Weeks 7–14).** Narrowest valuable slice: **Classroom + Calendar deadline radar + daily WhatsApp/PWA digest + NL Q&A**, sensitive-scopes only, "In Production" but unverified, <100 users, home campus. **Gate:** ⩾40% W4 retention, ⩾30% digest open rate.
- **Phase 2 — Home-campus pilot & iterate (Weeks 15–24).** Complete OAuth verification (sensitive scopes), push to a few hundred users, add per-college circular parsing. **Gate:** retention holds with verified app + >100 users; ⩾1 unsolicited "can my friends at [other college] use this?"
- **Phase 3 — Freemium launch + 2nd/3rd campus (Months 7–10).** Ship premium tier (Gmail + multimodal + risk score; accept CASA cost now). Launch ambassador program. **Gate:** ⩾1% conversion OR a signed college/department pilot.
- **Phase 4 — Multi-campus scale + funding (Months 11–18).** Template per-college onboarding; pursue institutional licensing; raise from college incubator/E-cell, government student-startup schemes, or angels using retention + B2B2C traction.

Gate for raising: demonstrable multi-campus retention + a repeatable per-college acquisition cost.

9. Risks & Honest Pitfalls

- **Google deprecates/expands access or ships parity in Gemini.** *Mitigation:* don't depend on Gmail early (use Classroom/Calendar); own the per-college data Google won't touch; build community lock-in.
- **OAuth/CASA friction kills momentum.** *Mitigation:* the sensitive-scopes-first path above; budget time for verification; keep Gmail (restricted) for later/premium.
- **Low willingness to pay.** *Mitigation:* lead with B2B2C/institutional revenue; keep per-user cost near-zero so free tier is sustainable; annual pricing.
- **Exam-season churn (students leave after exams).** *Mitigation:* the product is most valuable *during* the semester; make it a year-round habit (timetable, internship deadlines, club events), not just exam reminders.
- **Accuracy/hallucination erodes trust on deadlines (fatal for this category).** *Mitigation:* never silently auto-create events from low-confidence extractions; always show source + a one-tap confirm; treat a wrong deadline as a sev-1 bug. Trust is the entire product.
- **Privacy backlash / DPDP non-compliance.** *Mitigation:* DPDP Act 2023 + DPDP Rules 2025 (notified 13 Nov 2025) apply, with full compliance required by 13 May 2027. You are a Data Fiduciary processing personal data: implement explicit, granular, withdrawable consent; a plain-language privacy notice; data minimization; 72-hour breach-reporting readiness; and note that users under 18 require *verifiable parental consent* — so target 18+ college students and verify age. The Act gives startups a lighter graded burden, and data-minimization (sensitive scopes, on-device where possible) doubles as your DPDP and Google-trust story. Vakil Search
- **Solo-founder burnout.** *Mitigation:* narrow scope (one feature, one campus); automate ops; recruit a co-founder or first ambassador-turned-teammate before multi-campus scale.

10. The Narrow Wedge — The One Thing to Build First

Build exactly this: a "What's due this week?" deadline radar for one campus, powered by Google Classroom + Calendar (not Gmail), delivered as a daily WhatsApp/PWA digest with natural-language Q&A.

Why this and nothing else first:

- **Immediately lovable:** answers the one question every student asks weekly, with zero manual entry — the differentiator from My Study Life/Todoist.
- **Cheap to build:** Classroom + Calendar APIs, Llama-8B extraction, pgvector Q&A. ~₹2–4/user/month.
- **Compliance-light:** sensitive scopes only → no CASA, and a <100-user pilot needs no verification at all.
- **Defensible seed:** every correction trains your private deadline-extraction dataset; every campus adds per-college parsing depth.
- **Expandable:** Gmail, multimodal circulars/rubrics, the risk score, GitHub/Telegram all become *premium upgrades* later — not launch requirements.

Resist building the burnout/"academic risk score" first: it's the most speculative feature, the hardest to make accurate, and the easiest to erode trust with. Earn the right to predict overload by first nailing the boring, beloved deadline radar.

Recommendations (Prioritized "What To Do Next")

1. **This week:** Drop "Parsec/Singularity." Write the one-sentence wedge ("Never miss a deadline — your class assignments and college events, in one daily digest"). Stand up a Carrd/Vercel waitlist page.
2. **Weeks 1–6 (validation gate):** Run 15–20 Mom-Test interviews + a 10-person manual WhatsApp concierge digest at IIT Dharwad. **Only proceed if ⩾40% would be "very disappointed" without it.**

3. **Weeks 7–14:** Build the MVP on **Classroom + Calendar only** (sensitive scopes), "In Production" but unverified, <100 users. Stack: Python + Supabase (pgvector) + Groq Llama-8B + Gemini Flash-Lite + Vercel PWA + WhatsApp digest. Skip Pathway.
4. **Weeks 15–24:** Complete OAuth verification for sensitive scopes; iterate to ≥40% W4 retention; add per-college circular parsing. Apply for Microsoft/Cloudflare/Google startup credits now.
5. **Months 7–10:** Launch freemium (premium = Gmail + multimodal + risk score; accept ~$540/yr CASA only now). Start a campus ambassador program; pitch your placement cell/department for a paid or reference pilot.
6. **Months 11–18:** Template per-college onboarding; pursue B2B2C institutional licensing as the primary revenue; raise from your college E-cell/incubator or student-startup schemes on multi-campus retention.
7. **Always:** Treat a wrong deadline as a sev-1 trust bug; keep per-user cost near-zero; stay narrow and hyper-local; assume Google could enter and make per-college depth + community your moat.

Benchmarks that change the plan: If validation PMF <40% → pivot the wedge. If W4 retention <30% after MVP → fix retention before any growth spend. If free-to-paid <1% after freemium launch → pivot hard to B2B2C/institutional licensing and treat consumer as pure acquisition.

Caveats

- Several demand statistics (suicide figures, "61% of IIT attempts," US procrastination/email-overload rates) are macro/contextual and largely from secondary reporting; they establish stakes, not direct demand for *this* product. First-party campus validation is non-negotiable.
- OAuth scope classifications and the CASA "no-servers" exception are interpreted by Google reviewers case-by-case; treat the sensitive-vs-restricted distinction as the plan's foundation but verify current classification in your own Cloud Console before building, since Google has changed CASA processes mid-stream.

- LLM pricing and provider stability (Groq's NVIDIA transition, Gemini tier changes) shift frequently; keep a provider-abstraction layer.
- Unit-economics and conversion figures are order-of-magnitude estimates from public token pricing and typical usage/benchmark assumptions, not measured from a live cohort.
- DPDP Rules operational deadlines (e.g., consent-manager framework, full compliance by 13 May 2027) are still phasing in; monitor for the official Significant Data Fiduciary list and any education-sector guidance.